

Intro to Deep Learning Assignment 2 (Part 2)

Tsinghua University

Allan Li (李盼加乐), 2025403473

Part 1: Objective Questions

1. A
2. B
3. A
4. A
5. B, D
6. Two shortcomings of using an MLP are as follows:

1. First, an MLP treats a given image as a 1D vector of pixel values. The issue with this is that it doesn't capture the relationships of nearby pixels that might combine together to form a new edge/texture, because the model doesn't "know" which pixels might be adjacent to each other. Thus, it cannot naturally capture patterns like corners or shapes that might be core to the distinction of two different objects.
2. Secondly, in an image, each input pixel is connected to every neuron in the following layer. Thus, the number of parameters grow exponentially with image size. 10,000 pixels (100x100 image) connected to 100 hidden weights already needs 1 million weights. This might make MLPs unsuitable for image recognition because it is sensitive to overfitting, while also requiring a lot of computation power and time to train the model.

Part 2: Programming Practice

Results & console output can be found in `programming_practice.ipynb`

2.1

```
A = torch.rand(5, 5)
B = torch.rand(5, 5)
C = torch.zeros_like(A)

addition = A + B
subtraction = A - B
ele_multiplication = A * B
mat_multiplication = torch.matmul(A, B)
```

2.2

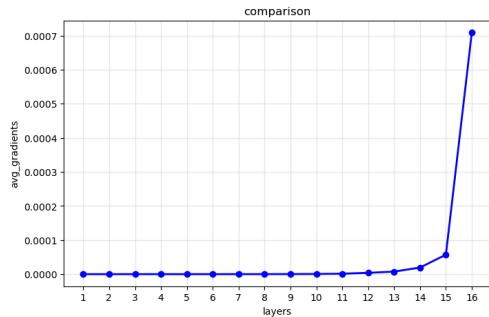
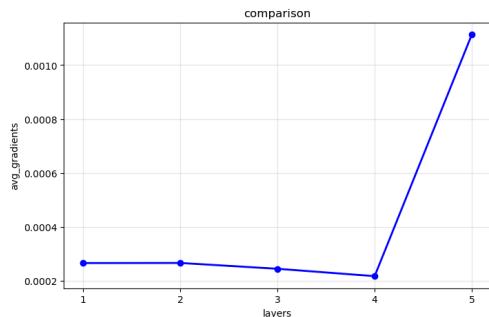
```
A = torch.rand(6, 4)
A_reshaped = A.reshape(24)
A_subset = A[1:5, 1:3]
```

2.3

```
A = torch.rand(3, 3, requires_grad=True)
B = torch.tensor([[1., 0., 2.],
                 [1., 1., 0.],
                 [0., 3., 1.]])
C = A * B
S = C.sum()
S.backward()
```

2.4: Observing Gradients

The following table summarizes differences in gradients:

Model Adjustment	Gradient Analysis	Gradient Visualization																																		
Base Model (No Changes)	Base Model.	 <table border="1"><caption>Data for Base Model Gradient Visualization</caption><thead><tr><th>Layers</th><th>Avg. Gradients</th></tr></thead><tbody><tr><td>1</td><td>0.0000</td></tr><tr><td>2</td><td>0.0000</td></tr><tr><td>3</td><td>0.0000</td></tr><tr><td>4</td><td>0.0000</td></tr><tr><td>5</td><td>0.0000</td></tr><tr><td>6</td><td>0.0000</td></tr><tr><td>7</td><td>0.0000</td></tr><tr><td>8</td><td>0.0000</td></tr><tr><td>9</td><td>0.0000</td></tr><tr><td>10</td><td>0.0000</td></tr><tr><td>11</td><td>0.0000</td></tr><tr><td>12</td><td>0.0000</td></tr><tr><td>13</td><td>0.0000</td></tr><tr><td>14</td><td>0.0000</td></tr><tr><td>15</td><td>0.0000</td></tr><tr><td>16</td><td>0.0007</td></tr></tbody></table>	Layers	Avg. Gradients	1	0.0000	2	0.0000	3	0.0000	4	0.0000	5	0.0000	6	0.0000	7	0.0000	8	0.0000	9	0.0000	10	0.0000	11	0.0000	12	0.0000	13	0.0000	14	0.0000	15	0.0000	16	0.0007
Layers	Avg. Gradients																																			
1	0.0000																																			
2	0.0000																																			
3	0.0000																																			
4	0.0000																																			
5	0.0000																																			
6	0.0000																																			
7	0.0000																																			
8	0.0000																																			
9	0.0000																																			
10	0.0000																																			
11	0.0000																																			
12	0.0000																																			
13	0.0000																																			
14	0.0000																																			
15	0.0000																																			
16	0.0007																																			
Decrease Layer Count (to 4)	Gradients increase overall. With less layers each layer has a larger impact on reducing the overall loss	 <table border="1"><caption>Data for 4-Layer Model Gradient Visualization</caption><thead><tr><th>Layers</th><th>Avg. Gradients</th></tr></thead><tbody><tr><td>1</td><td>0.0003</td></tr><tr><td>2</td><td>0.00025</td></tr><tr><td>3</td><td>0.0002</td></tr><tr><td>4</td><td>0.0002</td></tr><tr><td>5</td><td>0.0010</td></tr></tbody></table>	Layers	Avg. Gradients	1	0.0003	2	0.00025	3	0.0002	4	0.0002	5	0.0010																						
Layers	Avg. Gradients																																			
1	0.0003																																			
2	0.00025																																			
3	0.0002																																			
4	0.0002																																			
5	0.0010																																			

Model Adjustment	Gradient Analysis	Gradient Visualization																								
Increase Hidden Layer Count (to 512)	Wider layers cause the gradients to become relatively more uniform and stable (compared to base model)	<table border="1"> <caption>Data for Increase Hidden Layer Count</caption> <thead> <tr> <th>Layers</th> <th>Avg. Gradients</th> </tr> </thead> <tbody> <tr><td>1</td><td>0.00000</td></tr> <tr><td>2</td><td>0.00000</td></tr> <tr><td>3</td><td>0.00000</td></tr> <tr><td>4</td><td>0.00000</td></tr> <tr><td>5</td><td>0.00000</td></tr> <tr><td>6</td><td>0.00000</td></tr> <tr><td>7</td><td>0.00000</td></tr> <tr><td>8</td><td>0.00000</td></tr> <tr><td>9</td><td>0.00002</td></tr> <tr><td>10</td><td>0.00003</td></tr> <tr><td>11</td><td>0.00035</td></tr> </tbody> </table>	Layers	Avg. Gradients	1	0.00000	2	0.00000	3	0.00000	4	0.00000	5	0.00000	6	0.00000	7	0.00000	8	0.00000	9	0.00002	10	0.00003	11	0.00035
Layers	Avg. Gradients																									
1	0.00000																									
2	0.00000																									
3	0.00000																									
4	0.00000																									
5	0.00000																									
6	0.00000																									
7	0.00000																									
8	0.00000																									
9	0.00002																									
10	0.00003																									
11	0.00035																									
Decrease Learning Rate (to 0.0001)	Generally similar to the base model - somewhat more polarized gradients.	<table border="1"> <caption>Data for Decrease Learning Rate</caption> <thead> <tr> <th>Layers</th> <th>Avg. Gradients</th> </tr> </thead> <tbody> <tr><td>1</td><td>0.00000</td></tr> <tr><td>2</td><td>0.00000</td></tr> <tr><td>3</td><td>0.00000</td></tr> <tr><td>4</td><td>0.00000</td></tr> <tr><td>5</td><td>0.00000</td></tr> <tr><td>6</td><td>0.00000</td></tr> <tr><td>7</td><td>0.00000</td></tr> <tr><td>8</td><td>0.00000</td></tr> <tr><td>9</td><td>0.00002</td></tr> <tr><td>10</td><td>0.00005</td></tr> <tr><td>11</td><td>0.00080</td></tr> </tbody> </table>	Layers	Avg. Gradients	1	0.00000	2	0.00000	3	0.00000	4	0.00000	5	0.00000	6	0.00000	7	0.00000	8	0.00000	9	0.00002	10	0.00005	11	0.00080
Layers	Avg. Gradients																									
1	0.00000																									
2	0.00000																									
3	0.00000																									
4	0.00000																									
5	0.00000																									
6	0.00000																									
7	0.00000																									
8	0.00000																									
9	0.00002																									
10	0.00005																									
11	0.00080																									

Summary of results:

- Generally speaking, the gradients for all models were highest near the latter layers, with a massive increase in the last layer compared to all other layers. This makes sense since they are the last layer before the results, so changing them would heavily influence the results of the model
- Decreasing Layer count and increasing the hidden layer/making layers wider generally caused less polarization between the former and latter layers.
- Learning rate generally didn't change the gradient uniformity, but extreme values of learning rates on both sides would likely make gradients more polarized because the model instability increases

Simple Linear Regression Model

```

torch.manual_seed(42)
N = 256
x = torch.linspace(-2, 2, N).unsqueeze(1)
true_m, true_b = 3.0, 0.7
y = true_m * x + true_b + 0.2 * torch.randn_like(x)

class SimpleLinear(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc = nn.Linear(1, 1)
    
```

```

    def forward(self, x):
        return self.fc(x)

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = SimpleLinear().to(device)
criterion = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.05)

model.train()
data, target = x.to(device), y.to(device)

optimizer.zero_grad()
output = model(data)
loss = criterion(output, target)
loss.backward()
optimizer.step()

m_learned = model.fc.weight.item()
b_learned = model.fc.bias.item()
print(f"MSE Training loss after 1 epoch: {loss.item()}")
print(f"Learned params: w={m_learned}, b={b_learned}")

save_path = "weights.pt"
torch.save(model.state_dict(), save_path)
print(f"Saved weights to: {save_path}")

new_model = SimpleLinear().to(device)
new_model.load_state_dict(torch.load(save_path, map_location=device))
new_model.eval()

with torch.no_grad():
    test_x = torch.tensor([[-1.5], [0.0], [1.5]]).to(device)
    test_y_hat = new_model(test_x)
print(f"True params: w={true_m}, b={true_b}")
print("Predictions for x=[-1.5, 0.0, 1.5]:",
    test_y_hat.squeeze().cpu().numpy())
print(f"Actual y values: {true_m * test_x.squeeze().cpu().numpy() + true_b}")

```

Results:

- MSE Training loss after 1 epoch: 10.019256591796875
- Learned params: w=0.8872767686843872, b=-0.5553232431411743
- True params: w=3.0, b=0.7
- Predictions for x=[-1.5, 0.0, 1.5]: [-1.8862385 -0.55532324 0.77559197]

- Actual y values: [-3.8 0.7 5.2]